

CS4102 Algorithms

Spring 2020

Today's Keywords

- Greedy Algorithms
- Choice Function
- Cache Replacement
- Hardware & Algorithms

CLRS Reading: Chapter 16

Homeworks

- HW6 Due Sunday, April 5 @ 11pm
 - Written (use latex)
 - DP and Greedy
- HW10A also due Sunday, April 5 @ 11pm
 - No office hours offered

Caching Problem

- Why is using too much memory a bad thing?

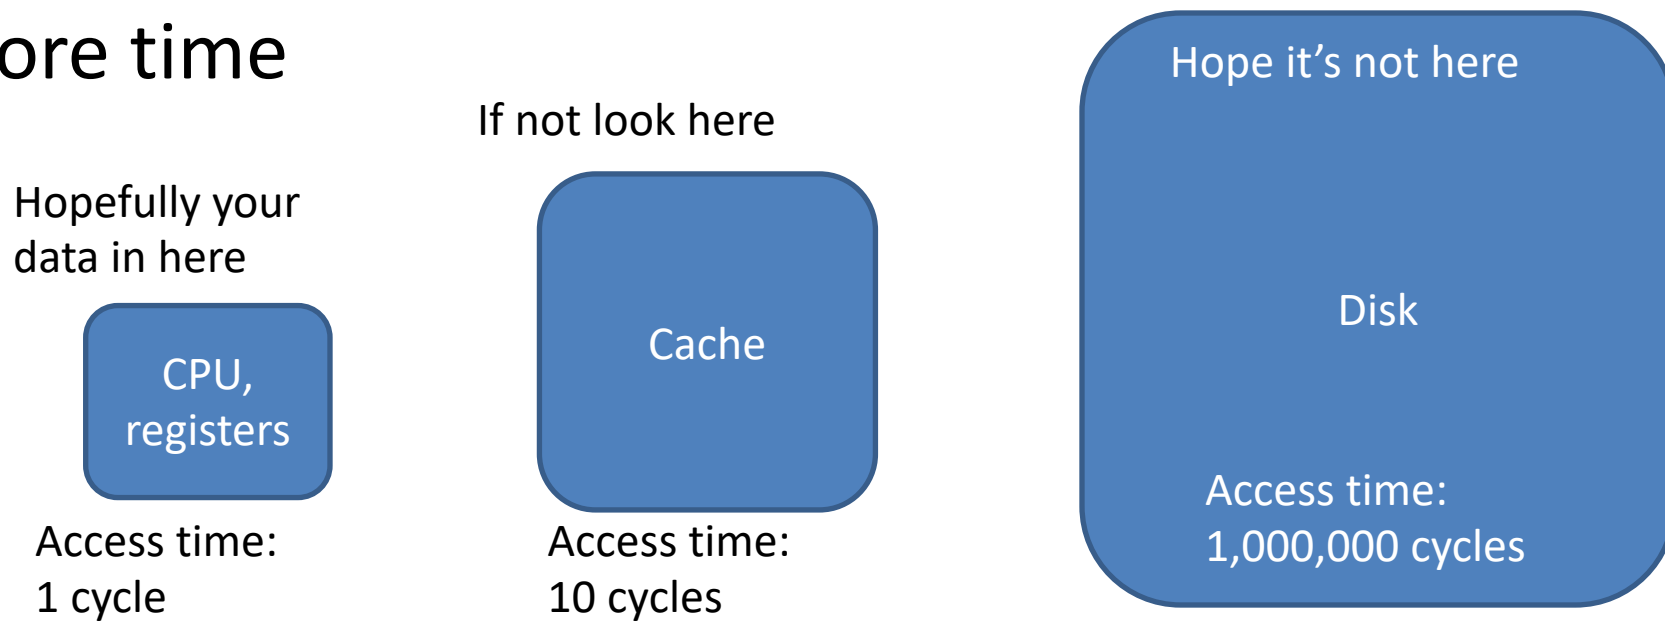
Von Neumann Bottleneck

- Named for John von Neumann
- Inventor of modern computer architecture
- Other notable influences include:
 - Mathematics
 - Physics
 - Economics
 - Computer Science



Von Neumann Bottleneck

- Reading from memory is VERY slow
- Big memory = slow memory
- Solution: hierarchical memory
- Takeaway for Algorithms: Memory is time, more memory is a lot more time



Caching Problem

- Cache misses are very expensive
- When we load something new into cache, we must eliminate something already there
- We want the best cache “schedule” to minimize the number of misses

Caching Problem Definition

- Input:
 - k = size of the cache
 - $M = [m_1, m_2, \dots, m_n]$ = memory access pattern
- Output:
 - “schedule” for the cache (list of items in the cache at each time)
which minimizes cache fetches

Example



A B C D A D E A D B A E C E A

Example



Example

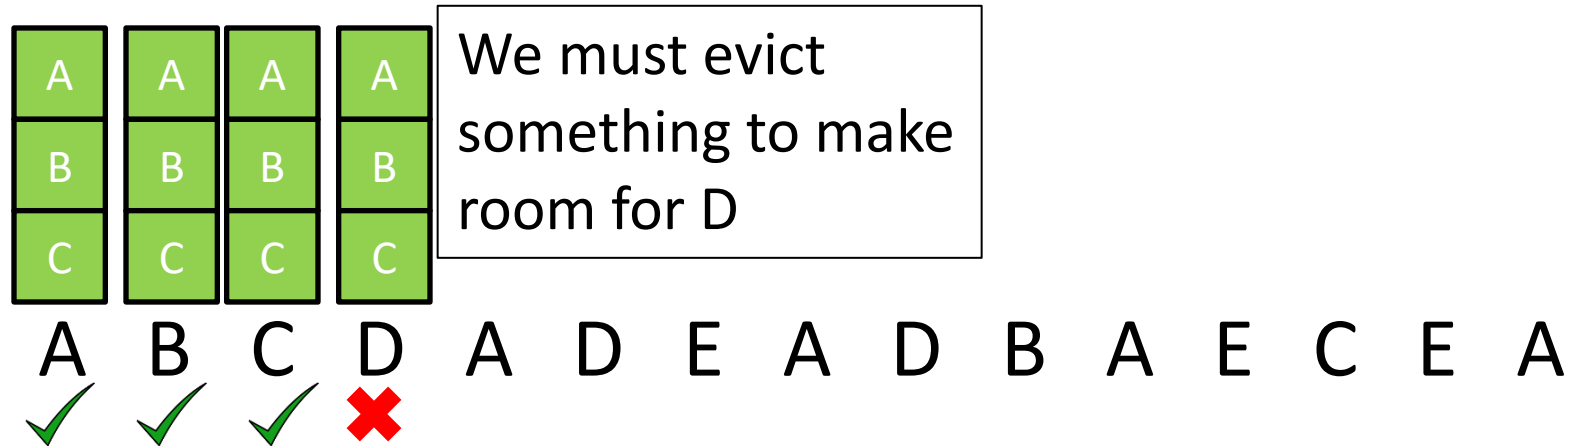
A	A	A
B	B	B
C	C	C

A B C

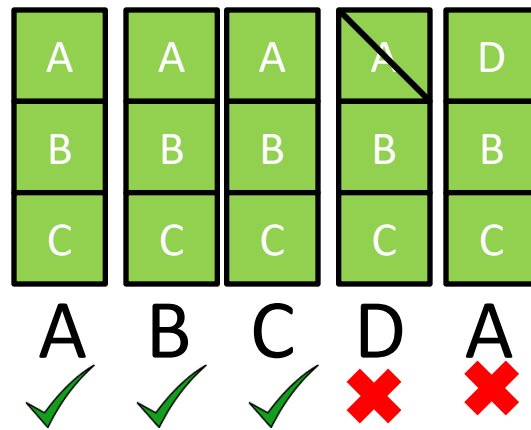
✓ ✓ ✓

D A D E A D B A E C E A

Example



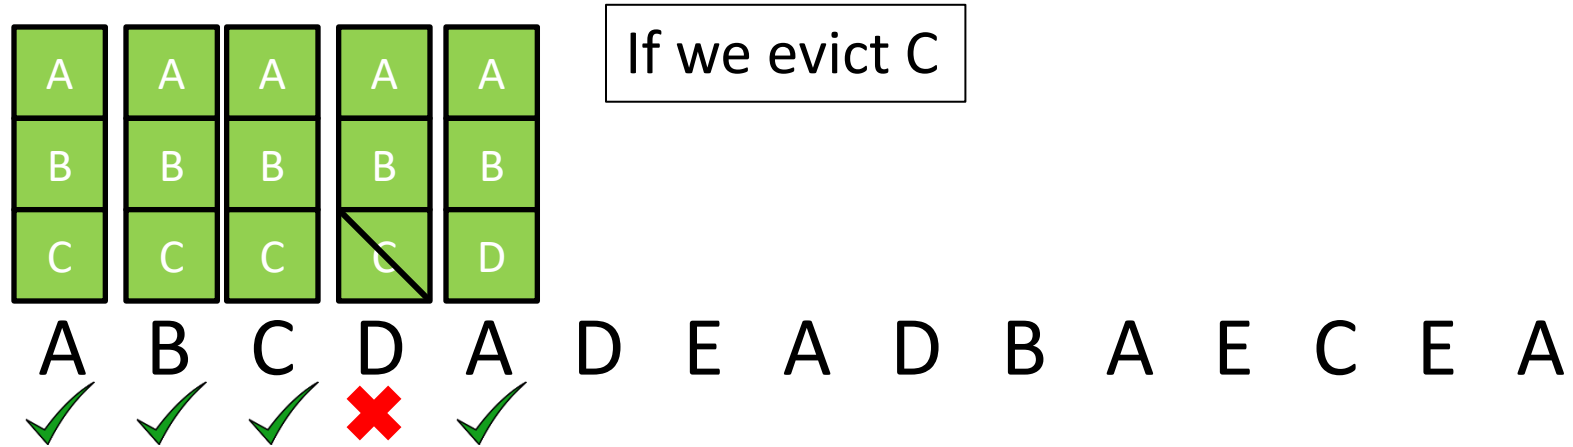
Example



If we evict A

D E A D B A E C E A

Example

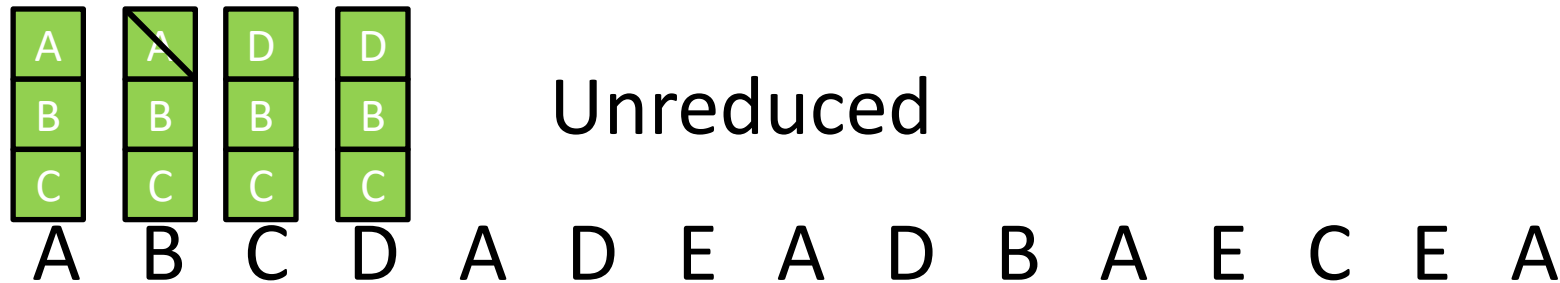


Our Problem vs Reality

- Assuming we know the entire access pattern
- Cache is Fully Associative
- Counting # of fetches (not necessarily misses)
- “Reduced” Schedule: Address only loaded on the cycle it’s required
 - Reduced == Unreduced (by number of fetches)

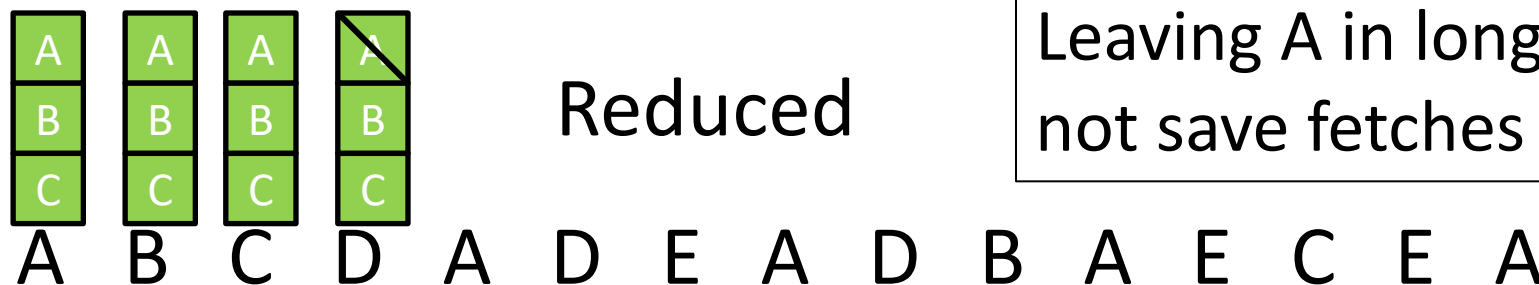
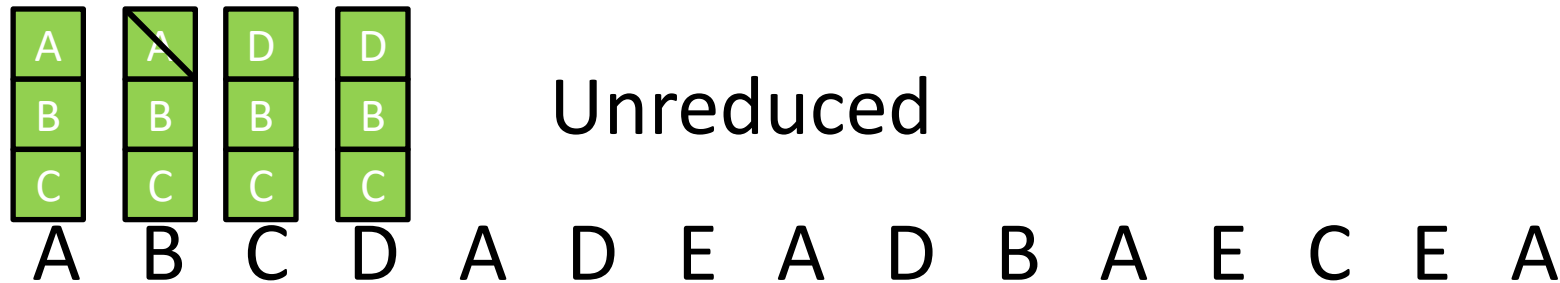
Our Problem vs Reality

- Assuming we know the entire access pattern
- Cache is Fully Associative
- Counting # of fetches (not necessarily misses)
- “Reduced” Schedule: Address only loaded on the cycle it’s required
 - Reduced == Unreduced (by number of fetches)



Our Problem vs Reality

- Assuming we know the entire access pattern
- Cache is Fully Associative
- Counting # of fetches (not necessarily misses)
- “Reduced” Schedule: Address only loaded on the cycle it’s required
 - Reduced == Unreduced (by number of fetches)



Leaving A in longer does not save fetches

Greedy Algorithms

- Require **Optimal Substructure**
 - Solution to larger problem contains the solution to a smaller one
 - Only one subproblem to consider!
- Idea:
 1. Identify a greedy **choice property**
 - How to make a choice guaranteed to be included in some optimal solution
 2. Repeatedly apply the choice property until no subproblems remain

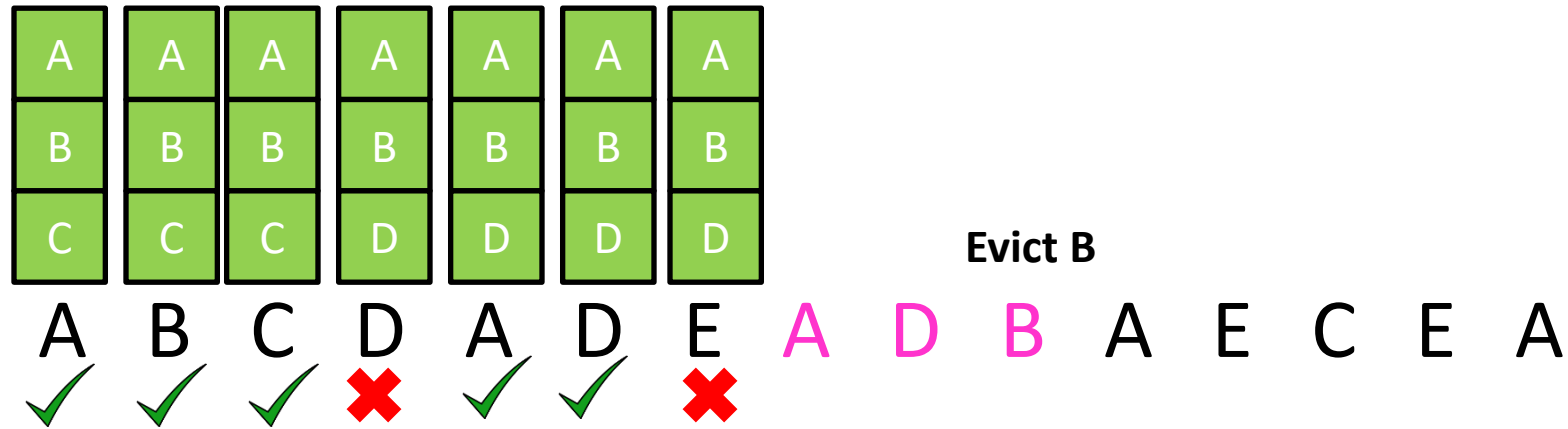
Greedy choice property

- Belady evict rule:
 - Evict the item accessed farthest in the future



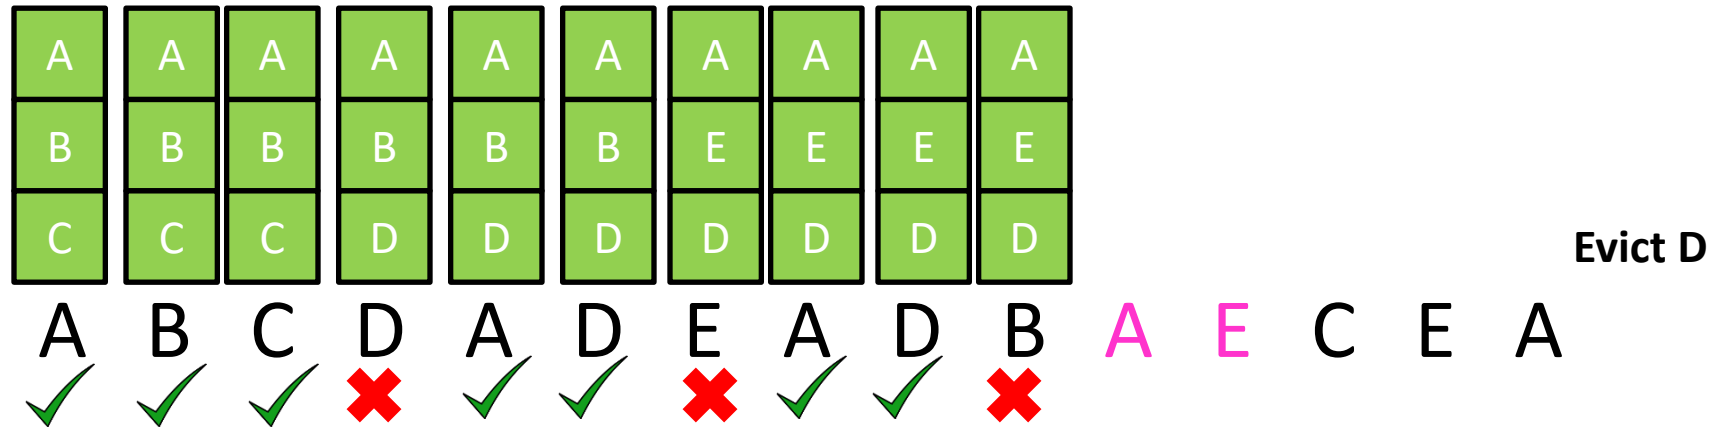
Greedy choice property

- Belady evict rule:
 - Evict the item accessed farthest in the future



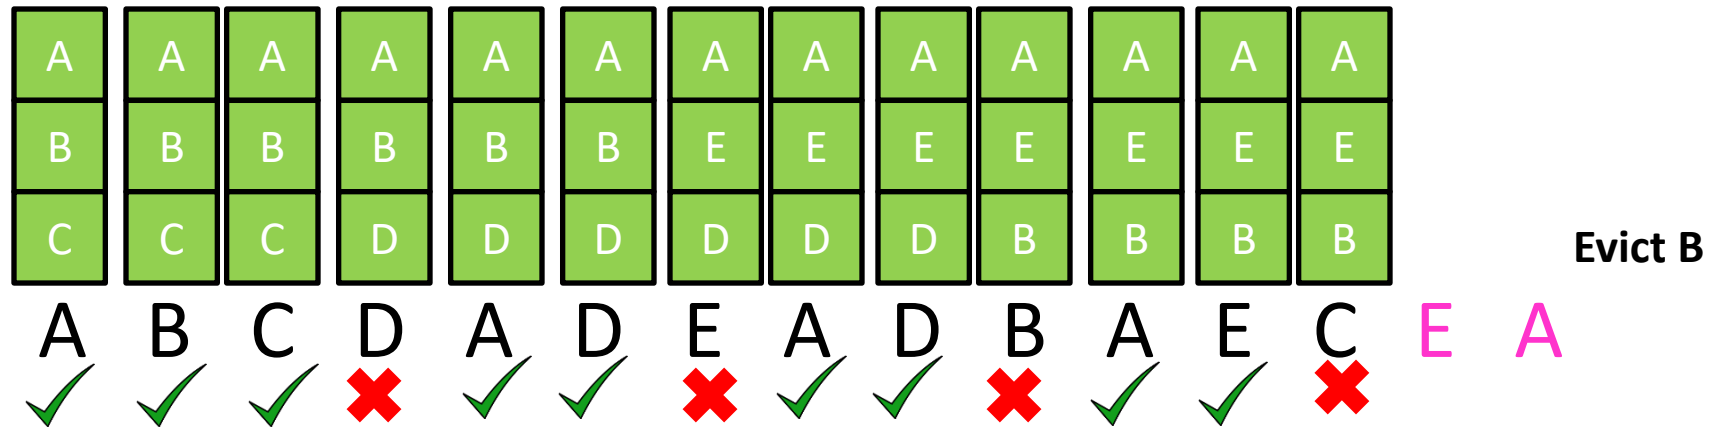
Greedy choice property

- Belady evict rule:
 - Evict the item accessed farthest in the future



Greedy choice property

- Belady evict rule:
 - Evict the item accessed farthest in the future



Greedy choice property

- Belady evict rule:
 - Evict the item accessed farthest in the future

A	A	A	A	A	A	A	A	A	A	A	A	A	A	A
B	B	B	B	B	B	E	E	E	E	E	E	E	E	E
C	C	C	D	D	D	D	D	D	B	B	B	C	C	C
A	B	C	D	A	D	E	A	D	B	A	E	C	E	A
✓	✓	✓	✗	✓	✓	✗	✓	✓	✗	✓	✓	✗	✓	✓

4 Cache Misses

Greedy Algorithms

- Require **Optimal Substructure**
 - Solution to larger problem contains the solution to a smaller one
 - Only one subproblem to consider!
- Idea:
 1. Identify a greedy **choice property**
 - How to make a choice guaranteed to be included in some optimal solution
 2. Repeatedly apply the choice property until no subproblems remain

Caching Greedy Algorithm

Initialize *cache* = first k accesses

For each $m_i \in M$:

 if $m_i \in \text{cache}$:

 print *cache*

 else:

m = furthest-in-future from cache

 evict m , load m_i

 print *cache*

Exchange argument

- Shows correctness of a greedy algorithm
- Idea:
 - Show exchanging an item from an arbitrary optimal solution with your greedy choice makes the new solution no worse
 - How to show my sandwich is at least as good as yours:
 - Show: “I can remove any item from your sandwich, and it would be no worse by replacing it with the same item from my sandwich”



Belady Exchange Lemma

Let S_{ff} be the schedule chosen by our greedy algorithm

Let S_i be a schedule which agrees with S_{ff} for the first i memory accesses.

We will show: there is a schedule S_{i+1} which agrees with S_{ff} for the first $i + 1$ memory accesses, and has no more misses than S_i
(i.e. $misses(S_{i+1}) \leq misses(S_i)$)

Belady Exchange Lemma

Let S_{ff} be the schedule chosen by our greedy algorithm

Let S_i be a schedule which agrees with S_{ff} for the first i memory accesses.

We will show: there is a schedule S_{i+1} which agrees with S_{ff} for the first $i + 1$ memory accesses, and has no more misses than S_i
(i.e. $misses(S_{i+1}) \leq misses(S_i)$)

Optimal



Agrees with
 S_{ff} on first 0
accesses

Belady Exchange Lemma

Let S_{ff} be the schedule chosen by our greedy algorithm

Let S_i be a schedule which agrees with S_{ff} for the first i memory accesses.

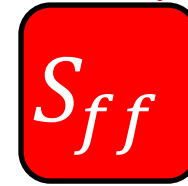
We will show: there is a schedule S_{i+1} which agrees with S_{ff} for the first $i + 1$ memory accesses, and has no more misses than S_i
(i.e. $misses(S_{i+1}) \leq misses(S_i)$)

Optimal



Agrees with
 S_{ff} on first 0
accesses

Greedy



Agrees with
 S_{ff} on all n
accesses

Belady Exchange Lemma

Let S_{ff} be the schedule chosen by our greedy algorithm

Let S_i be a schedule which agrees with S_{ff} for the first i memory accesses.

We will show: there is a schedule S_{i+1} which agrees with S_{ff} for the first $i + 1$ memory accesses, and has no more misses than S_i
(i.e. $misses(S_{i+1}) \leq misses(S_i)$)

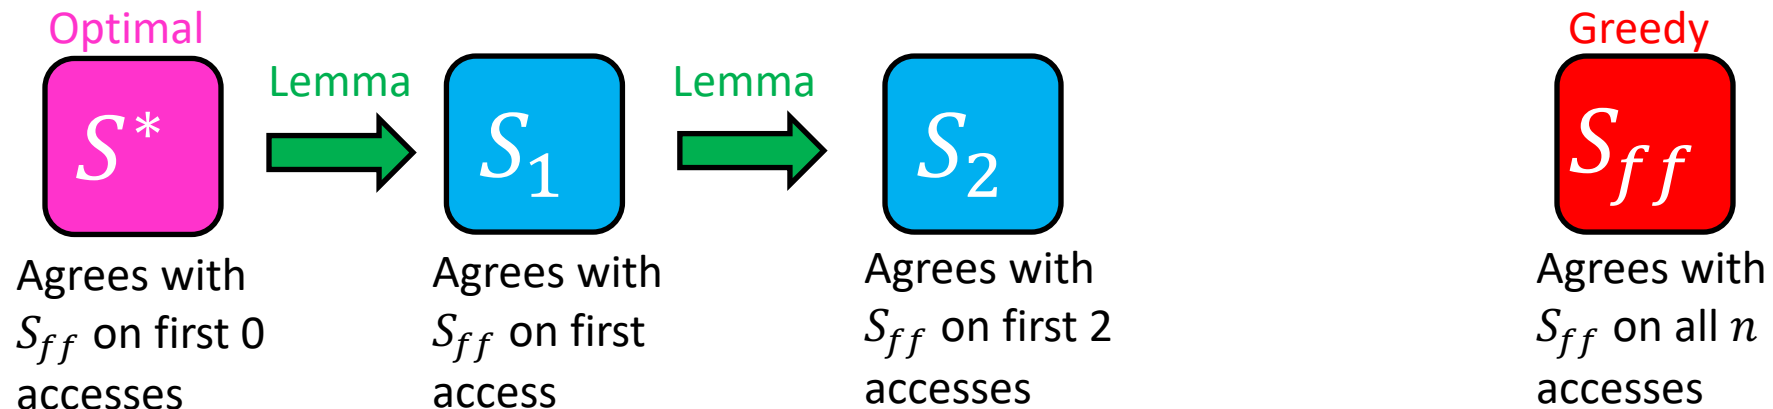


Belady Exchange Lemma

Let S_{ff} be the schedule chosen by our greedy algorithm

Let S_i be a schedule which agrees with S_{ff} for the first i memory accesses.

We will show: there is a schedule S_{i+1} which agrees with S_{ff} for the first $i + 1$ memory accesses, and has no more misses than S_i
(i.e. $misses(S_{i+1}) \leq misses(S_i)$)

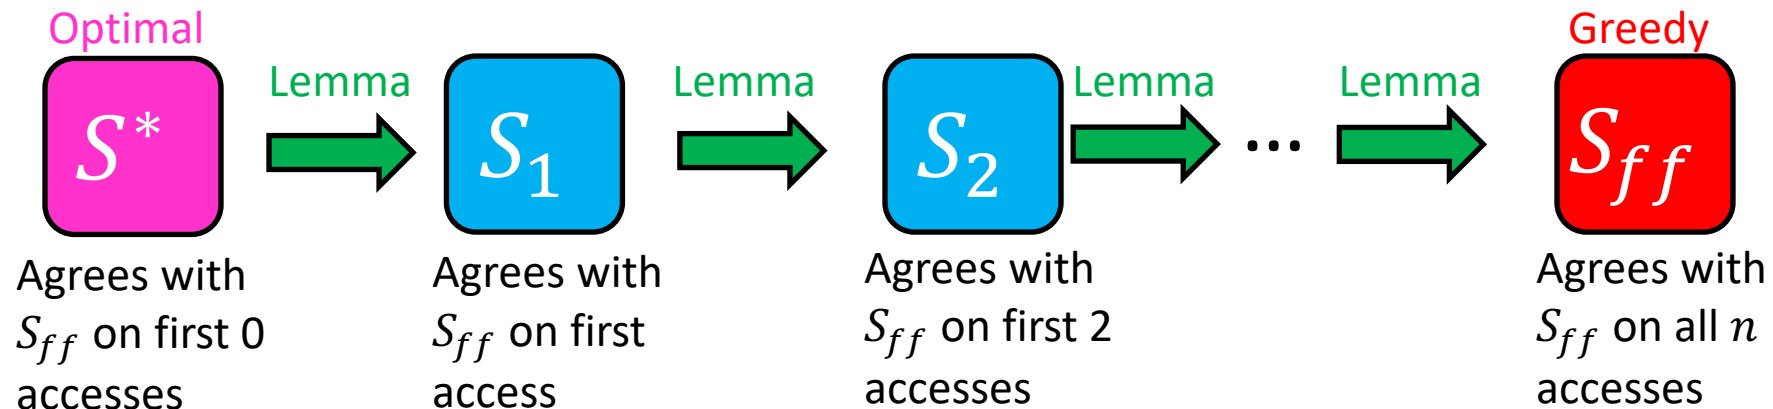


Belady Exchange Lemma

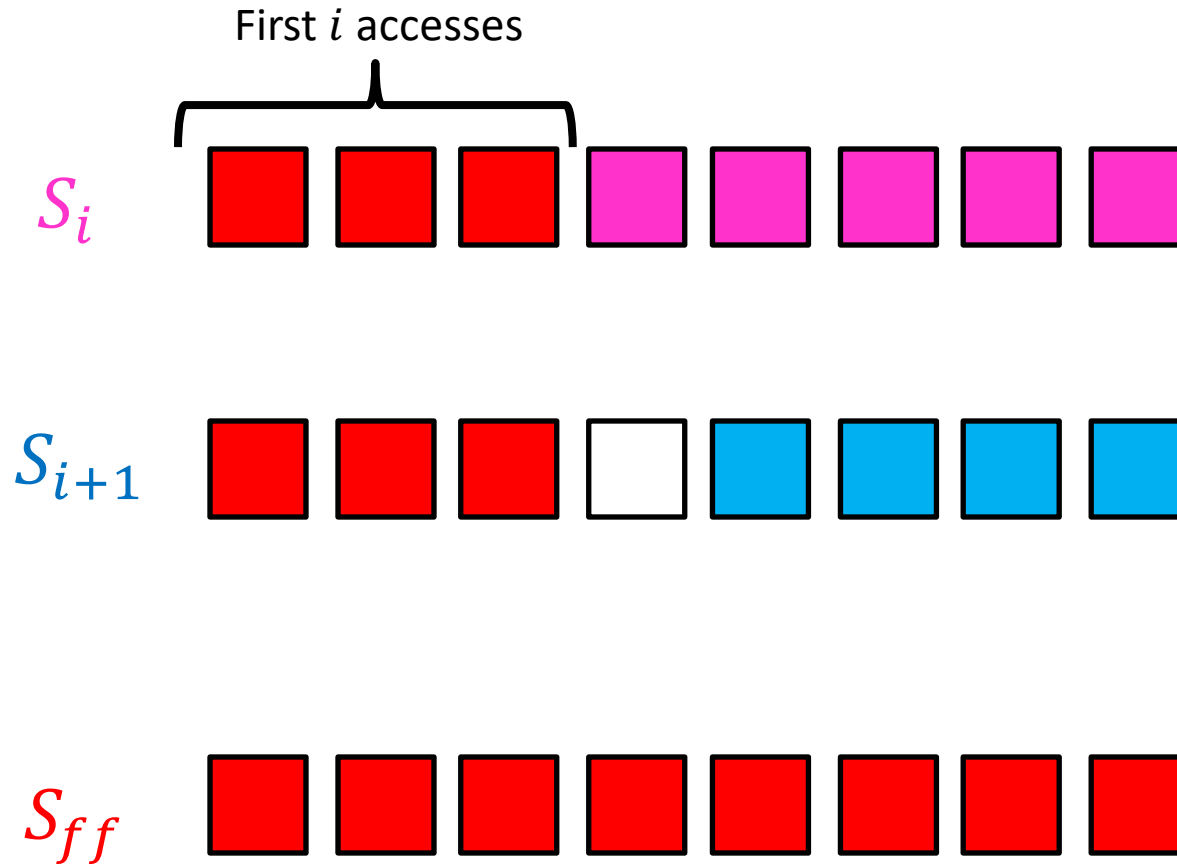
Let S_{ff} be the schedule chosen by our greedy algorithm

Let S_i be a schedule which agrees with S_{ff} for the first i memory accesses.

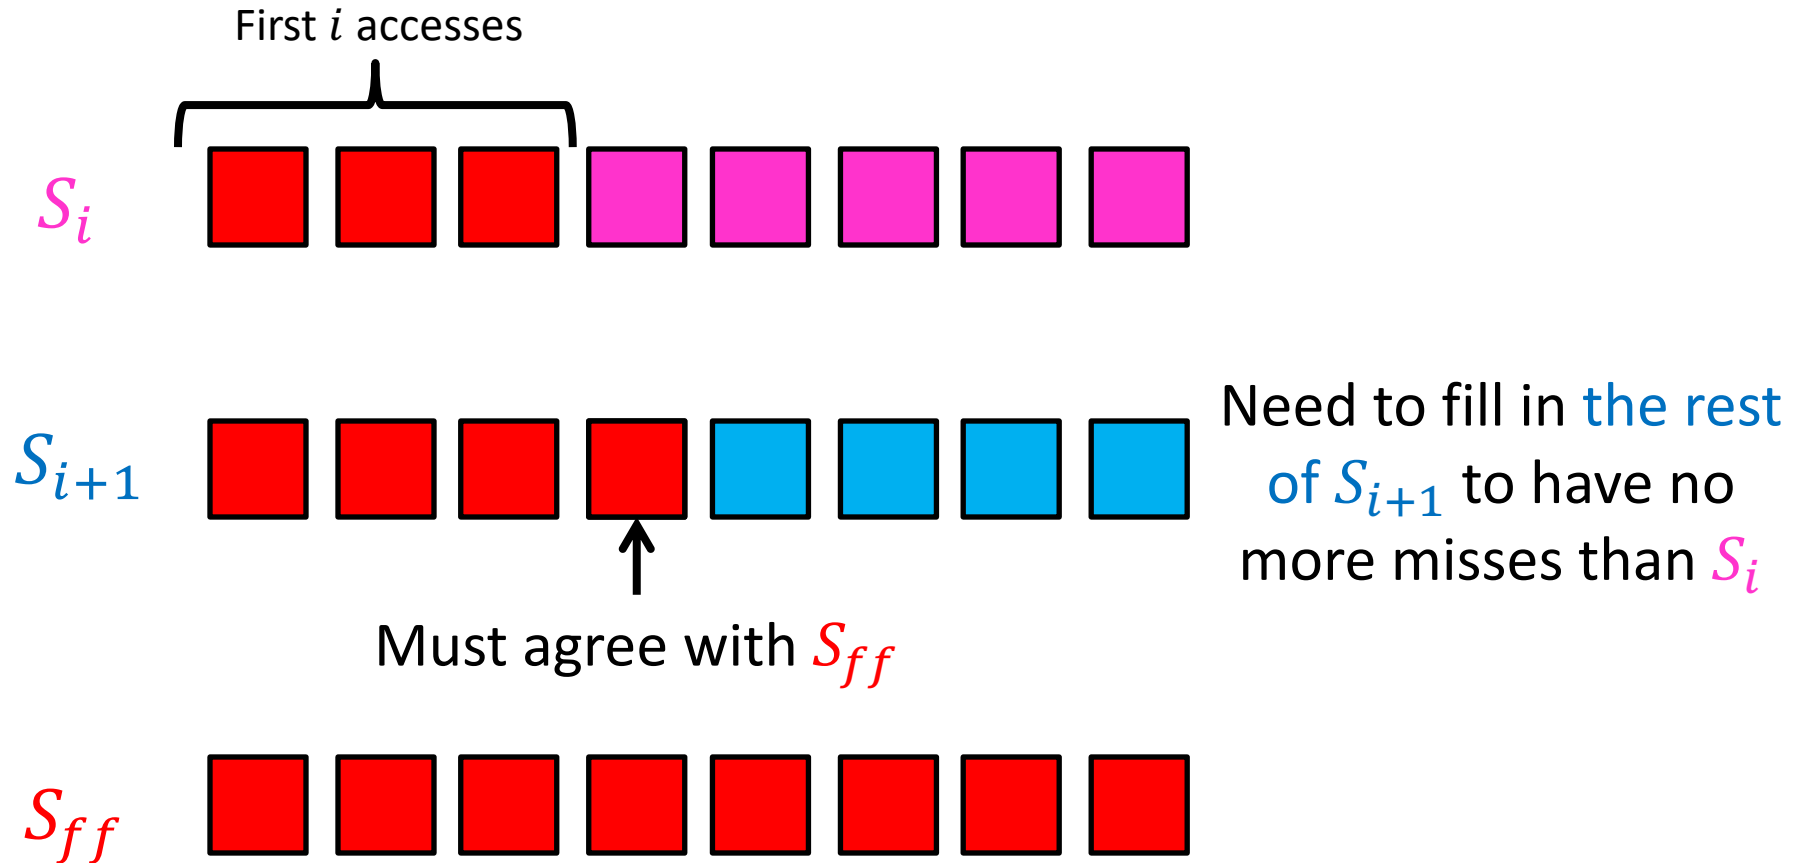
We will show: there is a schedule S_{i+1} which agrees with S_{ff} for the first $i + 1$ memory accesses, and has no more misses than S_i
(i.e. $misses(S_{i+1}) \leq misses(S_i)$)



Belady Exchange Proof Idea



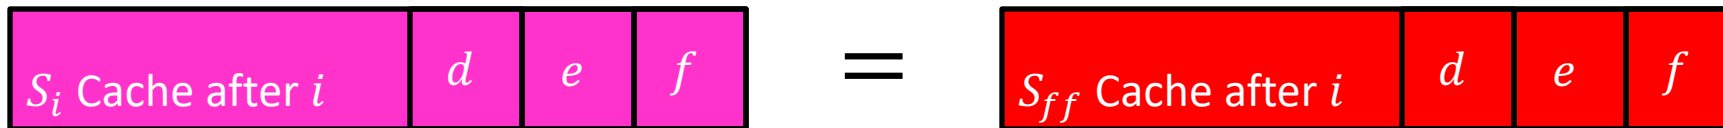
Belady Exchange Proof Idea



Proof of Lemma

Goal: find S_{i+1} s.t. $misses(S_{i+1}) \leq misses(S_i)$

Since S_i agrees with S_{ff} for the first i accesses, the state of the cache at access $i + 1$ will be the same



Consider access $m_{i+1} = d$

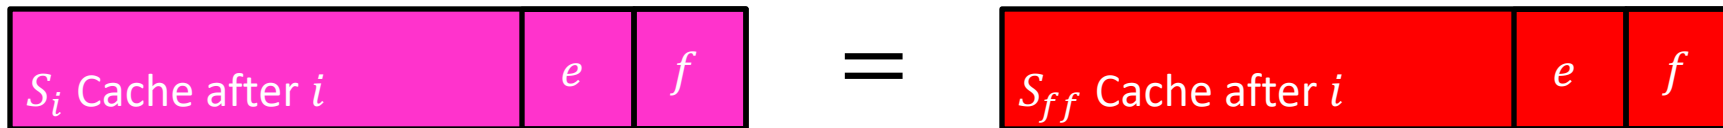
Case 1: if d is in the cache, then neither S_i nor S_{ff} evict from the cache, use the same cache for S_{i+1}



Proof of Lemma

Goal: find S_{i+1} s.t. $misses(S_{i+1}) \leq misses(S_i)$

Since S_i agrees with S_{ff} for the first i accesses, the state of the cache at access $i + 1$ will be the same



Consider access $m_{i+1} = d$

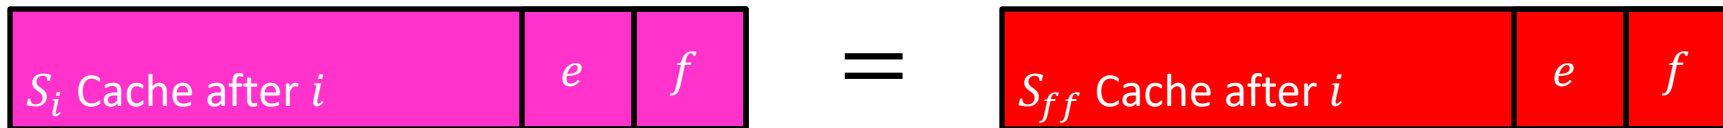
Case 2: if d isn't in the cache, and both S_i and S_{ff} evict f from the cache, evict f for d in S_{i+1}



Proof of Lemma

Goal: find S_{i+1} s.t. $\text{misses}(S_{i+1}) \leq \text{misses}(S_i)$

Since S_i agrees with S_{ff} for the first i accesses, the state of the cache at access $i + 1$ will be the same

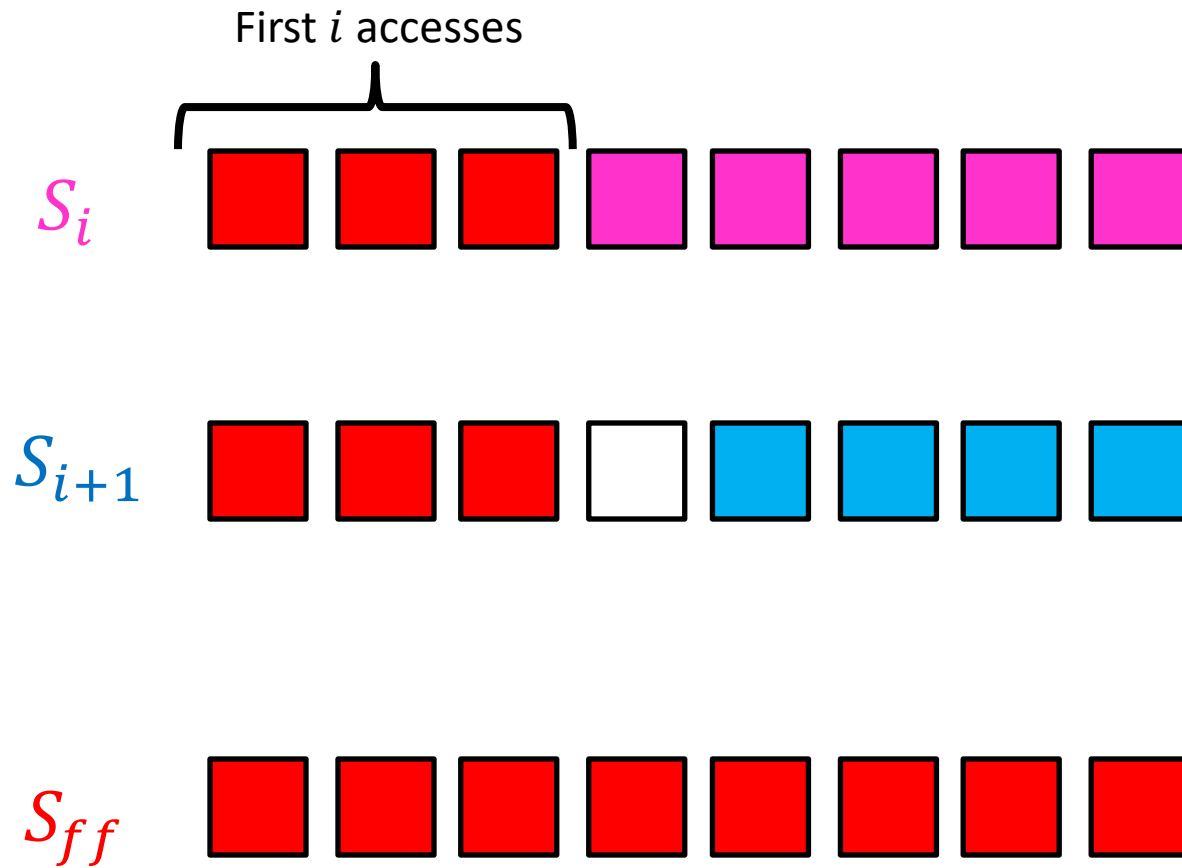


Consider access $m_{i+1} = d$

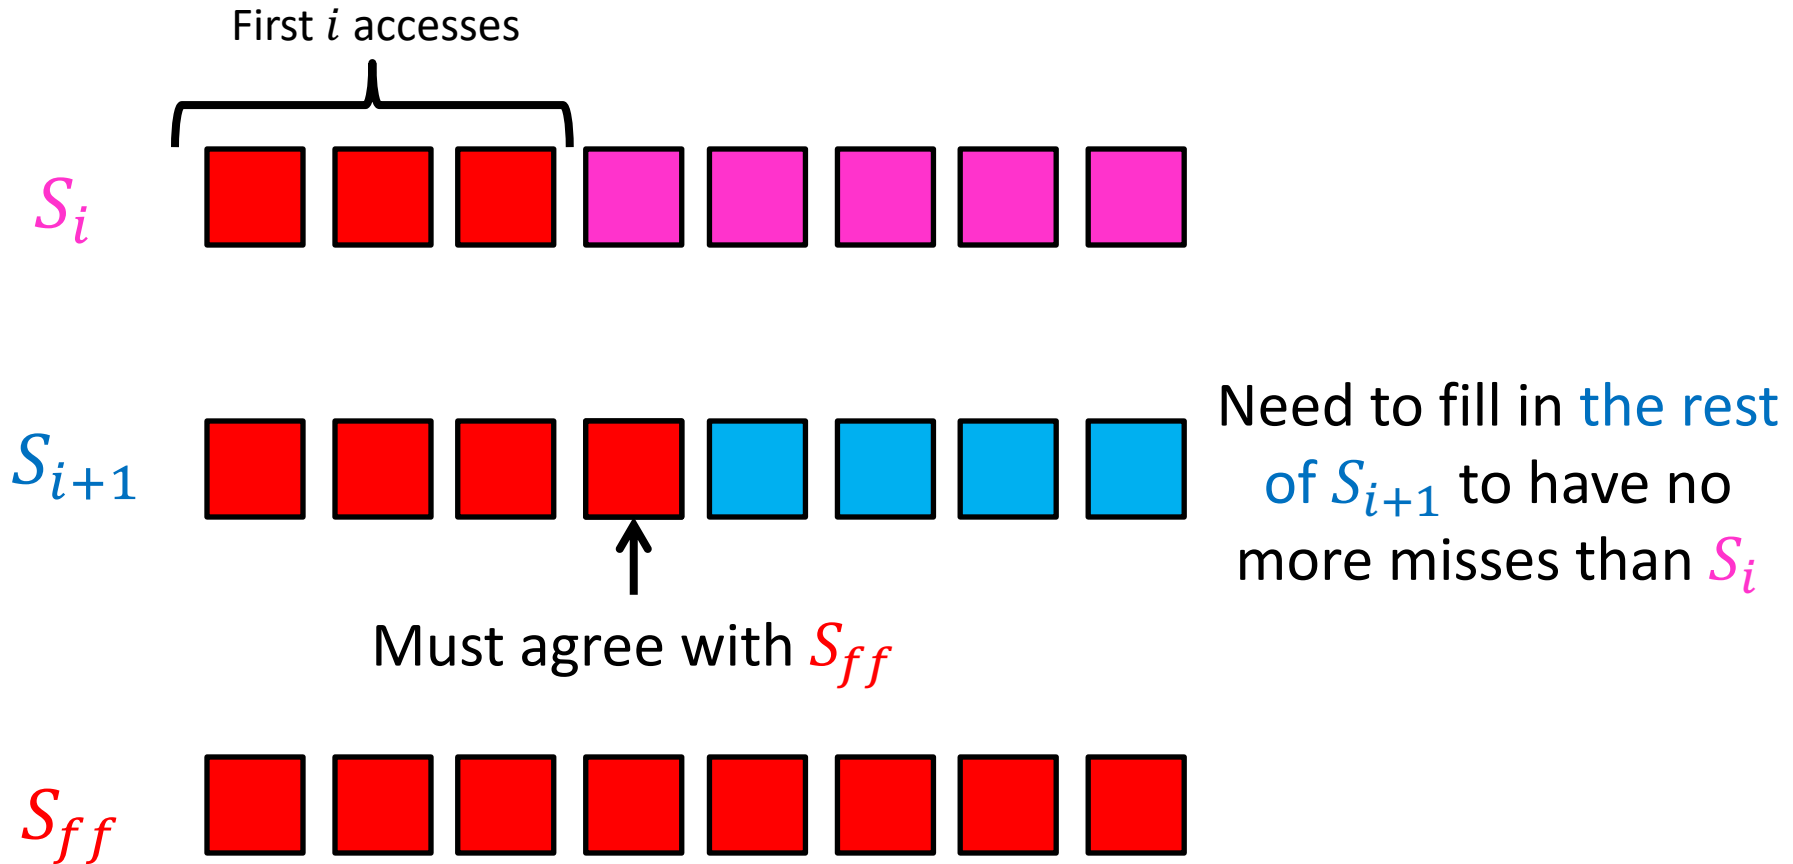
Case 3: if d isn't in the cache, S_i evicts e and S_{ff} evicts f from the cache



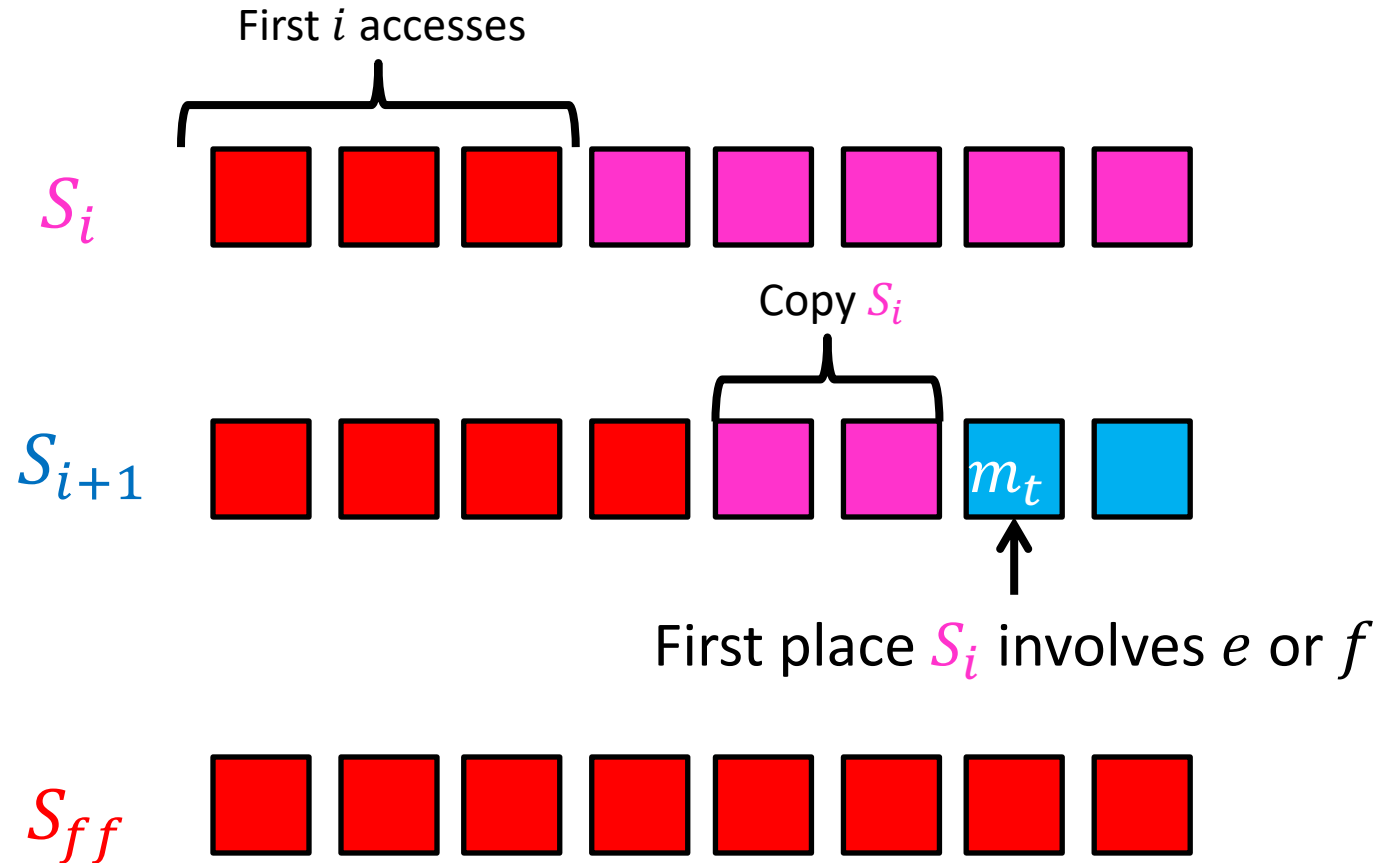
Case 3



Case 3



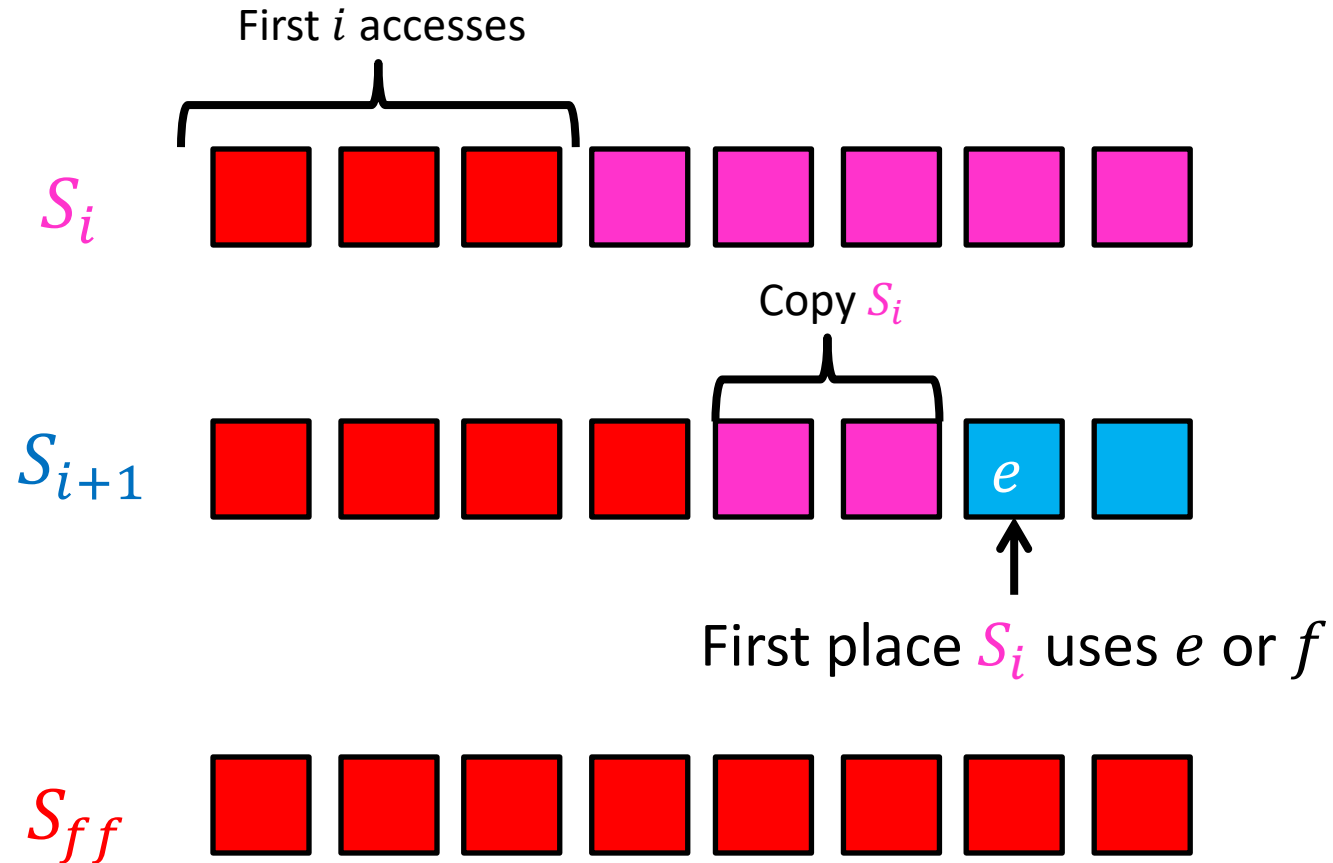
Case 3



m_t = the first access after $i + 1$ in which S_i deals with e or f

3 options: $m_t = e$ or $m_t = f$ or $m_t = x \neq e, f$

Case 3, $m_t = e$

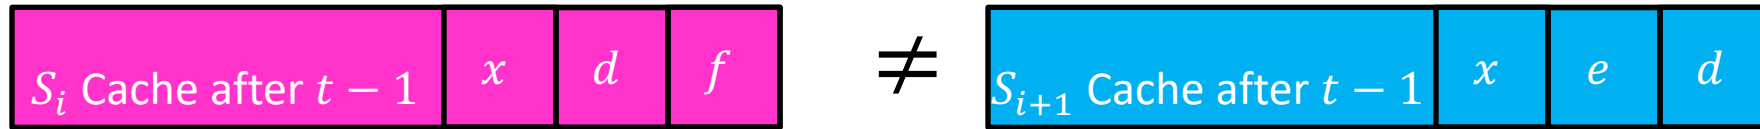


m_t = the first access after $i + 1$ in which S_i deals with e or f

3 options: $m_t = e$ or $m_t = f$ or $m_t = x \neq e, f$

Case 3, $m_t = e$

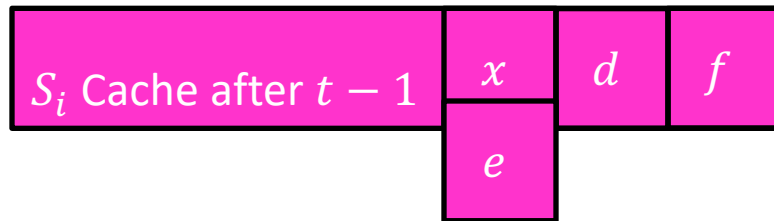
Goal: find S_{i+1} s.t. $\text{misses}(S_{i+1}) \leq \text{misses}(S_i)$



S_i must load e into
the cache, assume it
evicts x

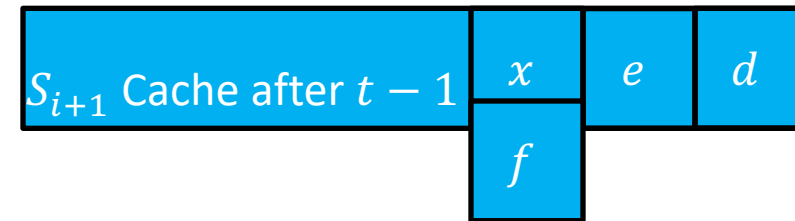
Case 3, $m_t = e$

Goal: find S_{i+1} s.t. $\text{misses}(S_{i+1}) \leq \text{misses}(S_i)$



S_i must load e into the cache, assume it evicts x

\neq

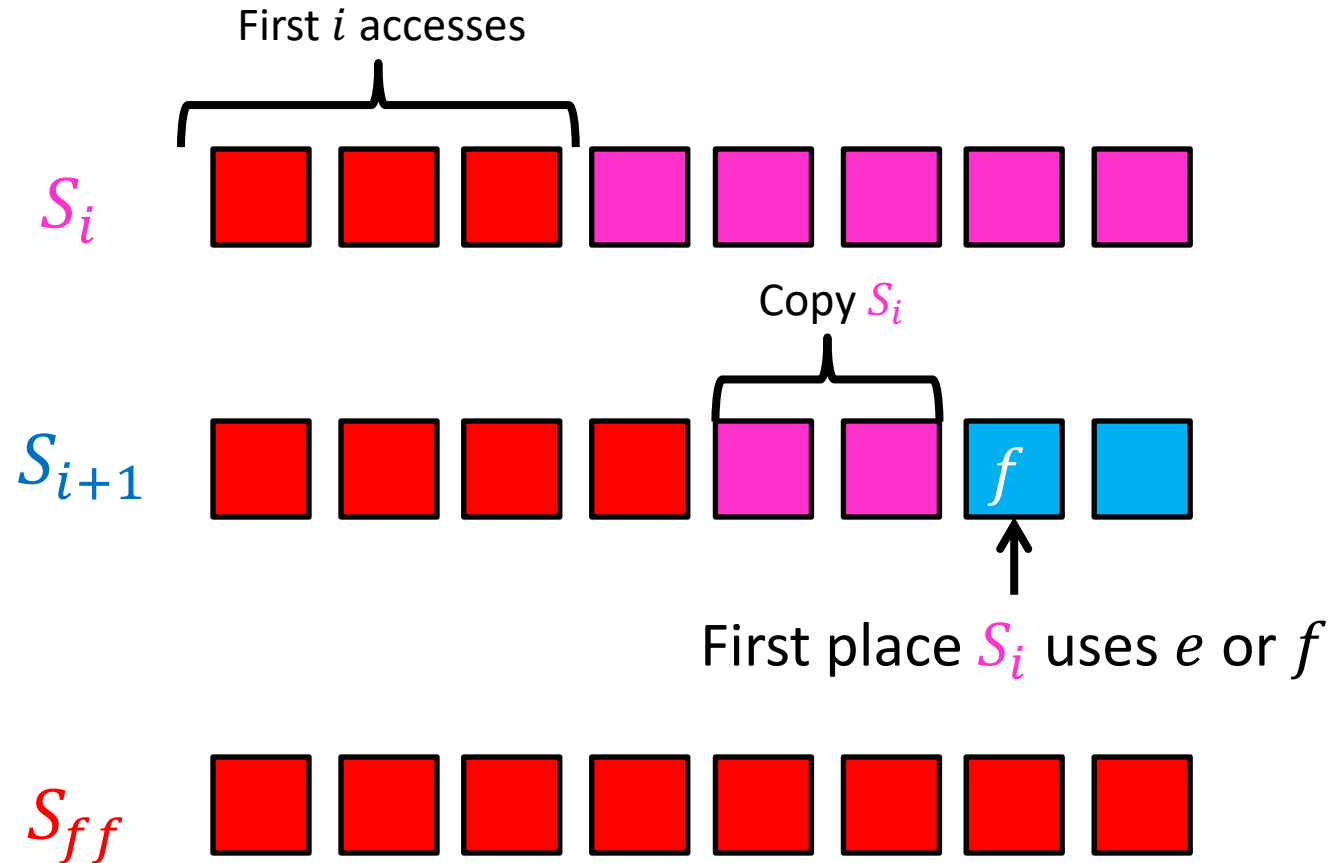


S_{i+1} will load f into the cache, evicting x

The caches now match!

S_{i+1} behaved exactly the same as S_i between i and t , and has the same cache after t , therefore $\text{misses}(S_{i+1}) = \text{misses}(S_i)$

Case 3, $m_t = f$

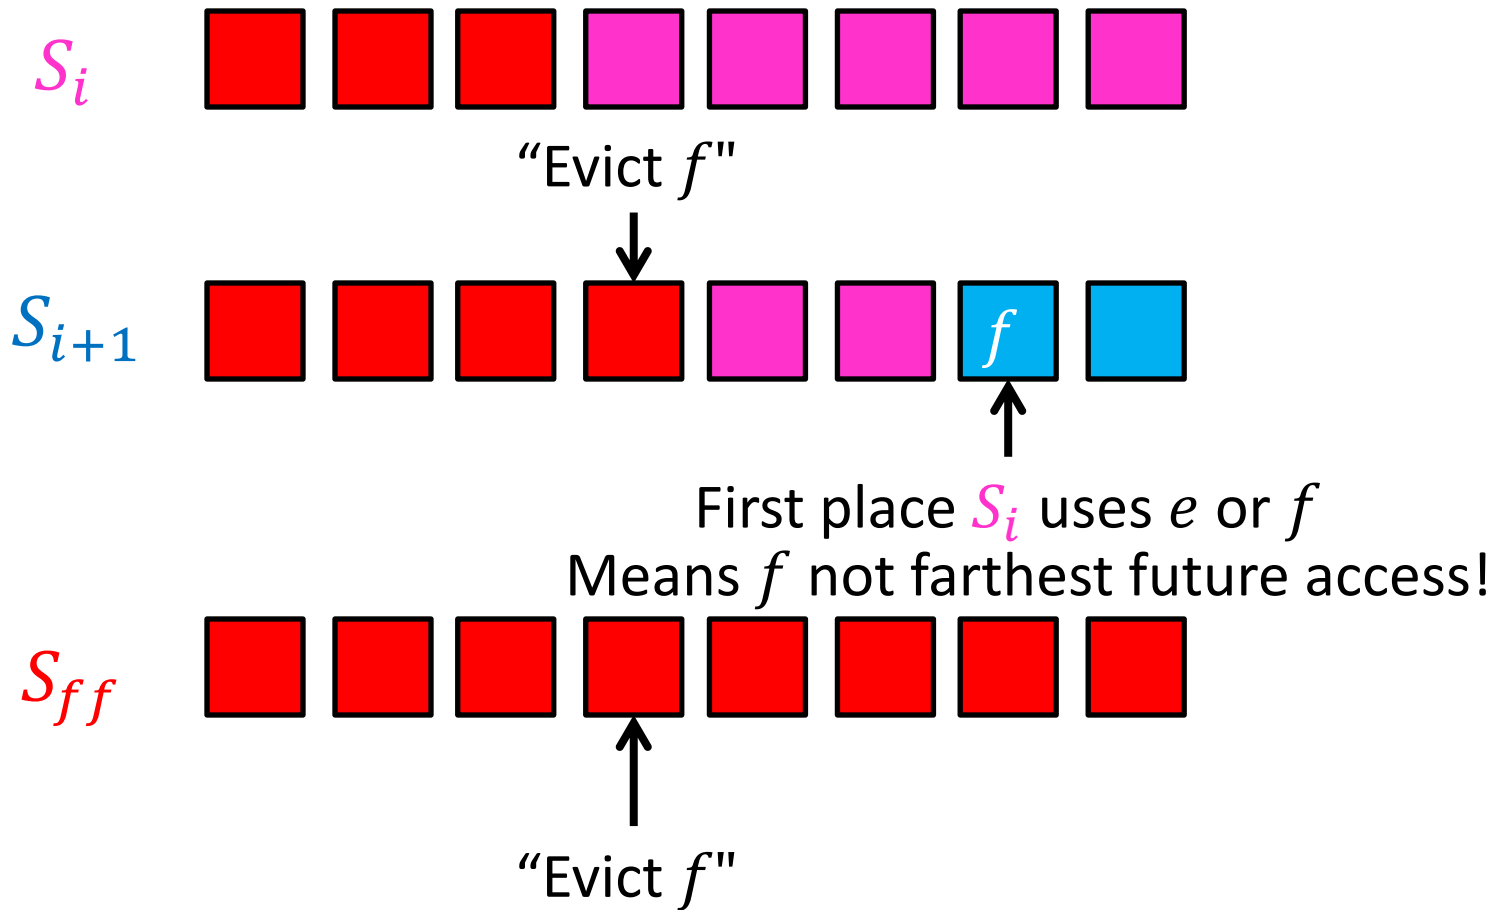


m_t = the first access after $i + 1$ in which S_i deals with e or f

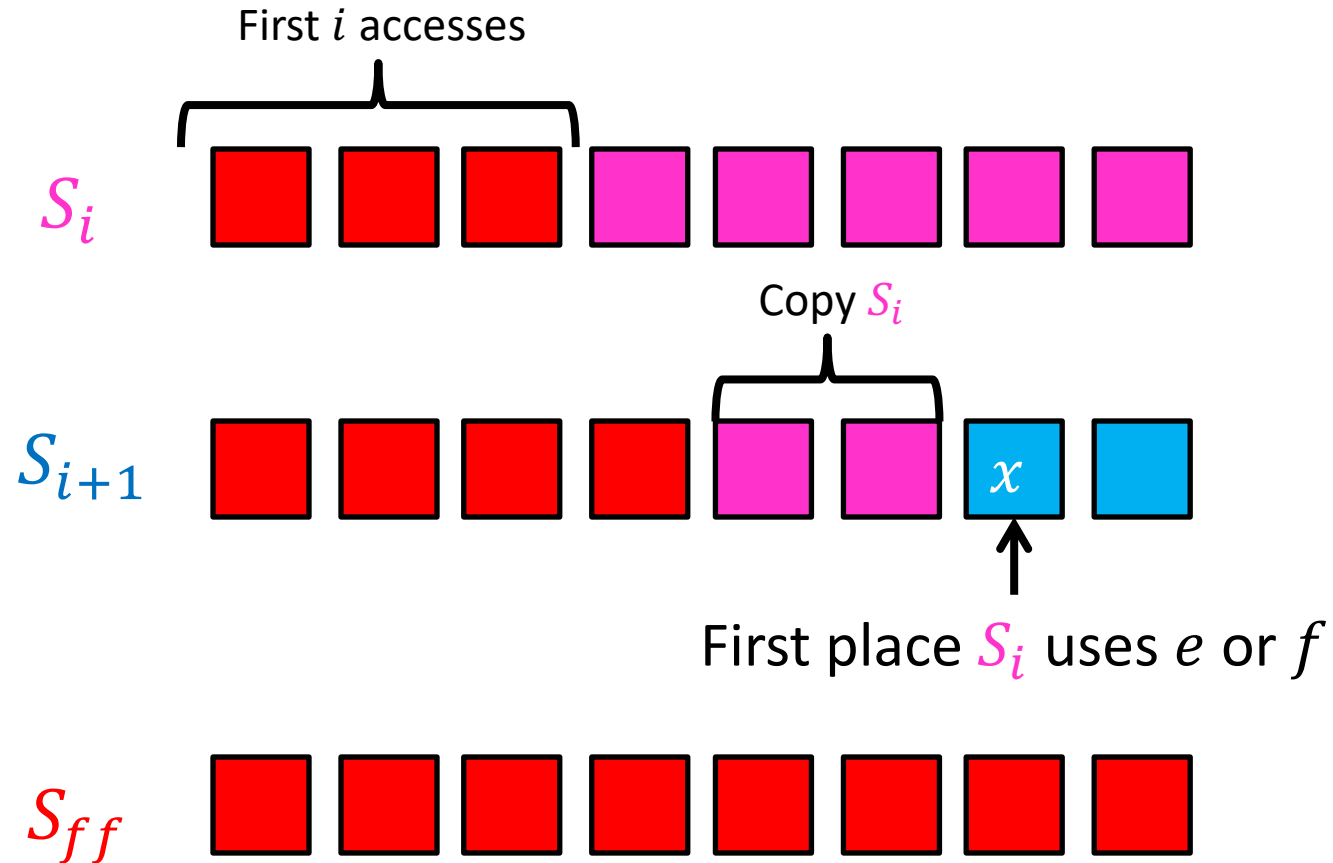
3 options: $m_t = e$ or $m_t = f$ or $m_t = x \neq e, f$

Case 3, $m_t = f$

Cannot Happen!



Case 3, $m_t = x \neq e, f$

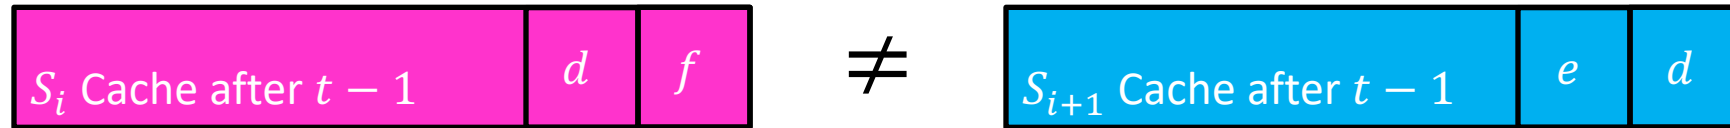


m_t = the first access after $i + 1$ in which S_i deals with e or f

3 options: $m_t = e$ or $m_t = f$ or $m_t = x \neq e, f$

Case 3, $m_t = x \neq e, f$

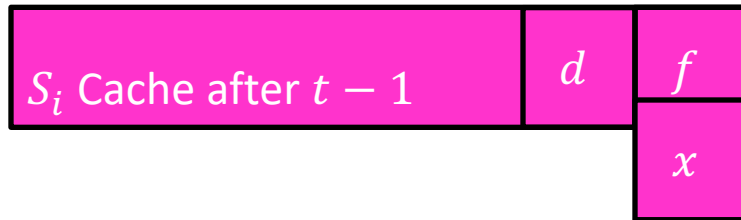
Goal: find S_{i+1} s.t. $\text{misses}(S_{i+1}) \leq \text{misses}(S_i)$



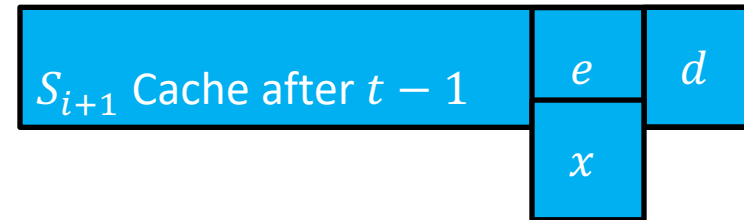
S_i loads x into the cache, it must be evicting f

Case 3, $m_t = x \neq e, f$

Goal: find S_{i+1} s.t. $\text{misses}(S_{i+1}) \leq \text{misses}(S_i)$



\neq



S_i loads x into the cache, it must be evicting f

S_{i+1} will load x into the cache, evicting e

The caches now match!

S_{i+1} behaved exactly the same as S_i between i and t , and has the same cache after t , therefore $\text{misses}(S_{i+1}) = \text{misses}(S_i)$

Use Lemma to show Optimality

